What is data mining? Explain data mining system architecture with necessary diagram.

1. What is Data Mining?

Data Mining is the process of extracting or "mining" useful knowledge, hidden patterns, and insights from massive amounts of raw data. It is the core step in the **Knowledge Discovery in Databases (KDD)** process.

In simple terms, just as mining for gold involves sifting through tons of earth to find a few precious nuggets, Data Mining sifts through huge databases to find valuable information that can be used for predicting trends, behavior analysis, and decision-making.

2. Data Mining System Architecture

A data mining system is composed of several integrated components that work together to transform raw data into knowledge. Below is the standard architecture:

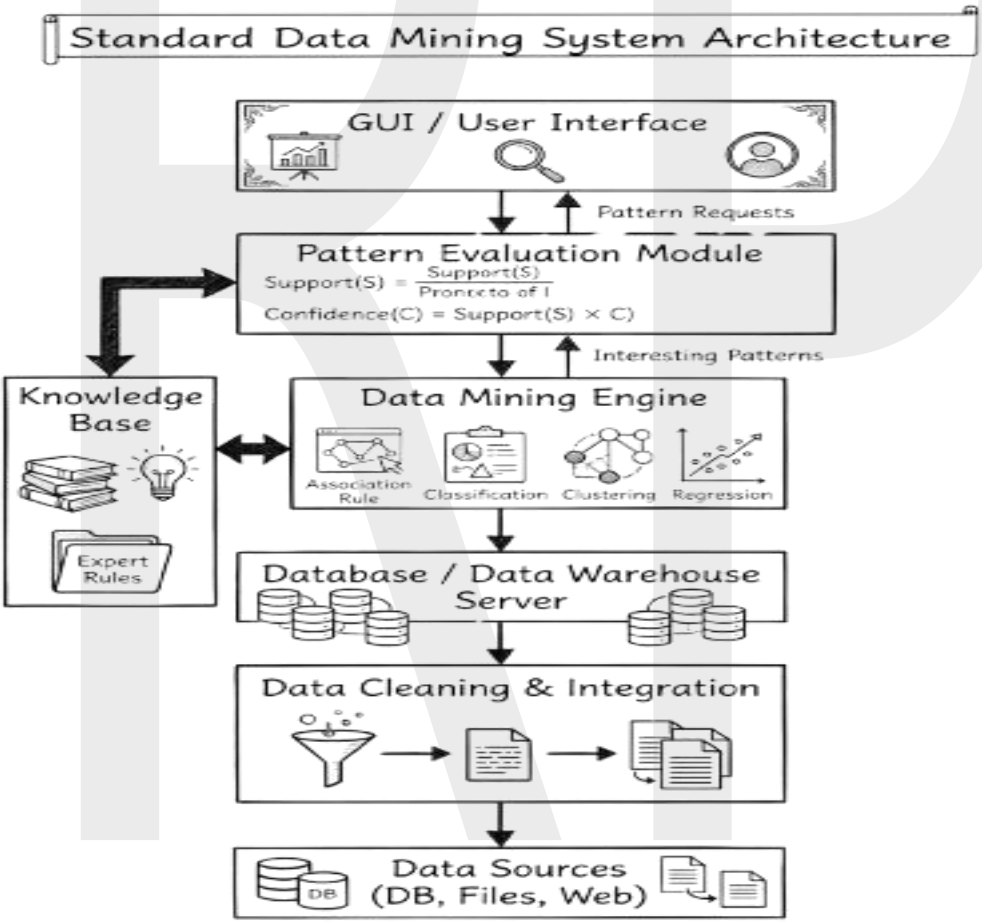

© Ramit Pariyar

Components of the Architecture:

- **Data Sources (Databases, Data Warehouses, WWW):** These are the repositories that hold the raw data. Data can be stored in relational databases, flat files, or gathered from the internet.
- **Database or Data Warehouse Server:** This component is responsible for fetching the relevant data based on the user's request. It performs basic operations like data retrieval and initial filtering.
- **Data Cleaning & Integration (Preprocessing):** Before mining, data must be cleaned (removing noise/errors) and integrated (combining data from multiple sources) to ensure accuracy.
- **Knowledge Base:** This is a critical component that stores **domain-specific knowledge**. It contains rules, constraints, and metadata that guide the search for patterns and help determine if a discovered pattern is actually useful or interesting.
- **Data Mining Engine:** This is the "heart" of the system. It consists of functional modules that perform tasks like **Association, Classification, Clustering, and Prediction**. It applies algorithms to the data to find patterns.
- **Pattern Evaluation Module:** This module works with the Mining Engine to filter out uninteresting patterns. It uses **interestingness measures** (like Support and Confidence) to ensure that only significant and useful patterns are presented to the user.
- **Graphical User Interface (GUI):** This is the top layer where the user interacts with the system. It allows the user to specify their mining tasks, visualize the results in the form of charts or tables, and explore the discovered knowledge.

Explain the prepocessing steps in detail.

Data Preprocessing is a crucial step in Data Mining because real-world data is often incomplete, noisy, and inconsistent. To get quality results from mining algorithms, we must first clean and prepare the data.

There are four major steps in Data Preprocessing:

© Ramit Pariyar

1. Data Cleaning

This step involves "cleaning" the data by filling in missing values and smoothing out noise.

- ✓ Handling Missing Values: If some data is missing, we can fill it with a global constant, the attribute mean, or the most probable value.
- ✓ Handling Noisy Data: Noise is a random error in a measured variable. We use Binning (sorting data into small "bins"), Regression, or Clustering to smooth the data and remove outliers.

2. Data Integration

Data often comes from multiple sources (different databases, files, or nodes). Data integration combines these sources into a single, unified data store.

- ✓ Entity Identification: Ensuring that "Student_ID" in one database and "S_ID" in another refer to the same entity.
- ✓ Redundancy Analysis: Removing duplicate data to save space and improve mining speed.

3. Data Transformation

In this step, data is converted or consolidated into forms appropriate for mining.

- ✓ Normalization: Scaling the data values to a specific range (e.g., 0.0 to $1.0).
- ✓ Aggregation: Summarizing data (e.g., daily sales data is aggregated to monthly totals).
- ✓ Generalization: Replacing low-level data with higher-level concepts (e.g., replacing "22 years old" with "Youth").

4. Data Reduction

Large datasets can take a long time to process. Data reduction obtains a reduced representation of the dataset that is much smaller in volume but produces the same (or almost the same) analytical results.

- ✓ Dimensionality Reduction: Removing irrelevant attributes (e.g., using Wavelet transforms or PCA).

© Ramit Pariyar

- ✓ Numerosity Reduction: Replacing the actual data with smaller forms like models or histograms.
- ✓ Data Compression: Encoding the data to reduce its size.

What re the characteristics of data warehouse? Differentiate between database and data warehouse.

1. Characteristics of a Data Warehouse

I. Subject-Oriented

Unlike operational databases that focus on processes (like sales, invoicing, or inventory), a data warehouse focuses on high-level **subjects** like Customers, Products, or Sales. It excludes data that is not useful for decision-making.

II. Integrated

Data is gathered from multiple heterogeneous sources (different databases, flat files, etc.). During the integration process, names, units, and structures are **standardized**.

Example: If one source uses "M/F" for gender and another uses "0/1", the data warehouse converts them all to a single format like "Male/Female".

III. Time-Variant

Data in a warehouse is identified with a particular **time period**. While a database usually shows only current data (like your current balance), a data warehouse stores historical data (like your balance over the last 10 years) to analyze trends.

IV. Non-Volatile

Once data is entered into the warehouse, it is **not changed or deleted**. It is permanent. The only operations performed are initial data loading and data access (querying). This ensures that reports generated today are consistent with reports generated in the future.

© Ramit Pariyar

2. Difference between Database and Data Warehouse

Feature	Database (OLTP)	Data Warehouse (OLAP)
Purpose	To record day-to-day transactions.	To analyze data for decision-making.
Data Type	Contains only **current** data.	Contains **historical** and current data.
Operations	Frequent Inserts, Updates, and Deletes.	Mostly **Read-only** (Bulk loading of data).
User	Clerks, IT Professionals.	Managers, Analysts, Executives.
Design	Optimized for fast transactions (Normalised).	Optimized for complex queries (Denormalised).
Size	Usually smaller (GBs).	Very large (TBs to PBs).

Explain about different OLAP operation.

Roll-up (Aggregation)

Roll-up performs aggregation on a data cube by either climbing up a concept hierarchy or by reducing dimensions. It moves from specific data to more general data.

Example: Moving from Daily sales data to Monthly sales data, or moving from City level to Country level.

Result: You get a broader overview but less detail.

2. Drill-down (Reverse of Roll-up)

Drill-down is the opposite of roll-up. It allows the user to move from summary data to more detailed data. This is done by moving down a concept hierarchy or adding a new dimension.

Example: Moving from Annual sales to Quarterly sales, or looking at sales for a specific Store instead of the whole Region.

Result: You get more specific details and a deeper look into the data.

3. Slice and Dice

These operations focus on selecting specific portions of the data cube.

- ✧ **Slice:** This operation selects a **single dimension** from the cube to create a new sub-cube (like taking a single "slice" of bread from a loaf).

Example: Selecting sales data for only the year "2025" from a cube containing many years.

- ✧ **Dice:** This operation selects a **sub-cube** by picking specific values from multiple dimensions.

Example: Selecting sales data for "Laptops" (Product) in "Nepal" (Location) during "March" (Time).

4. Pivot (Rotate)

Pivot, also known as **Rotation**, does not change the data but changes the **orientation** of the view. It rotates the data axes to provide an alternative presentation of the data.

- ✓ **Example:** Swapping the rows and columns in a spreadsheet. Instead of having "Products" in rows and "Cities" in columns, you rotate it so "Cities" are in rows and "Products" are in columns.
- ✓ **Result:** Helps in looking at the same data from a different visual perspective.

Find the frequent item sets and generate association rules on the following transaction using Apriori Algorithm.

Transaction ID	Items
T1	I1, I2, I3
T2	I1, I2
T3	I1, I4, I5
T4	I4, I5
T5	I5, I3

© Ramit Pariyar

Transaction ID Items

T6 I1, I4, I5

- Minimum Support = 33.33% = 2
- Minimum Confidence = 60%

Given:

- Minimum Support = 33.33% = 2 transactions
- Minimum Confidence = 60%

Step 1: Frequent 1-itemsets (L1)

Item Support Count

I1 4
I2 2
I3 2
I4 3
I5 4

L1 = {I1, I2, I3, I4, I5}

Step 2: Frequent 2-itemsets (L2)

Itemset Support Count

I1, I2 2

I1, I4 2

I1, I5 2

I4, I5 3

L2 = {(I1,I2), (I1,I4), (I1,I5), (I4,I5)}

Step 3: Frequent 3-itemsets (L3)

Itemset Support Count

I1, I4, I5 2

© Ramit Pariyar

L3 = {(I1,I4,I5)}

Step 4: Association Rules

Formula:
Confidence(A → B) = Support(A ∪ B) / Support(A)

Strong Association Rules (≥ 60%)

- I2 → I1 = 100%
- I4 → I5 = 100%
- I5 → I4 = 75%
- (I1, I4) → I5 = 100%
- (I1, I5) → I4 = 100%

Final Result:

Frequent Itemsets:

- L1 = {I1, I2, I3, I4, I5}
- L2 = {(I1,I2), (I1,I4), (I1,I5), (I4,I5)}
- L3 = {(I1,I4,I5)}

Strong Rules:

- I2 → I1
- I4 → I5
- I5 → I4
- (I1,I4) → I5
- (I1,I5) → I4

What is classification? Explain Bayesian classification with suitable example.

Classification

Classification is a data mining task that predicts the category (class label) of a given data point. It is a form of **Supervised Learning**, meaning the system is first trained on a labeled dataset (where the answer is known) before it is used to categorize new, unseen data.

Goal: To build a model that can map input data to a specific discrete class.

Example: Identifying whether an email is "Spam" or "Not Spam," or whether a patient has a certain "Disease" or is "Healthy."

Bayesian Classification

Bayesian Classification is based on the Bayes' Theorem. It is a statistical classifier that predicts the probability that a given tuple belongs to a particular class.

The most common version is the Naïve Bayes Classifier, which assumes that the effect of an attribute value on a given class is independent of the values of other attributes (this is called class conditional independence).

The Formula:

The fundamental formula for Bayes' Theorem is:

$$P(H|E) = \frac{P(E|H).\ P(H)}{P(E)}$$

Where:

P(H|E) is the **Posterior Probability** (Probability that hypothesis H is true given evidence E).

P(H) is the **Prior Probability** of the hypothesis.

P(E|H) is the **Likelihood** (Probability of evidence E given H is true).

P(E) is the **Prior Probability** of the evidence.

Suitable Example: Predicting "Will Play Football?"

Suppose we have a dataset with features like **Weather** and we want to classify if a person will **Play** or **Not**.

Training Data:

| Weather | Play (Target) |

| :--- | :--- |

| Sunny | No |

| Overcast | Yes |

| Rainy | Yes |

| Sunny | Yes |

| Overcast | Yes |

Total Records: 5

Yes: 4/5 (0.8)

No: 1/5 (0.2)

Define clustering. Explain about agglomerative and divisive clustering.

Clustering is the process of grouping a set of objects in such a way that objects in the same group (called a cluster) are more similar to each other than to those in other groups. Unlike classification, clustering is Unsupervised Learning, meaning the data does not have pre-defined class labels. The system finds patterns on its own based on the features of the data.

Goal: To achieve high intra-cluster similarity (objects in the same cluster are very alike) and low inter-cluster similarity (different clusters are very distinct).

Example: Segmenting customers into groups based on their purchasing habits or grouping similar species of plants together.

Hierarchical Clustering

Hierarchical clustering creates a tree-like structure of clusters called a **Dendrogram**. There are two main approaches to building this hierarchy:

1. Agglomerative Clustering (Bottom-Up)

This is the most common type of hierarchical clustering. It starts with every single data point as an individual cluster and merges them step-by-step.

Process: 1. Start with n clusters (each point is its own cluster).

2. Find the two closest clusters and merge them into one.

3. Repeat the process until all points are merged into a single large cluster.

Analogy: Like individuals coming together to form small groups, which then join to form a large organization.

2. Divisive Clustering (Top-Down)

This is the opposite of agglomerative clustering. It starts with one giant cluster containing all data

Start with one sin points and splits it repeatedly.

Process:

- ✓ gle cluster containing all n data points.
- ✓ Split the cluster into two most dissimilar sub-clusters.
- ✓ Repeat the process recursively on the sub-clusters until every point is its own individual cluster.
- ✓ **Analogy:** Like a large country dividing into provinces, and those provinces further dividing into districts.

Compare classification with clustering. Describe text mining in details.

Classification vs. Clustering

Feature	**Classification**	**Clustering**
Type of Learning	Supervised Learning (has a teacher/label).	Unsupervised Learning (no labels).
Goal	To predict the target class for new data.	To find natural groupings in the data.
Data Labeling	Uses a Training Set with known labels.	Uses Unlabeled Data.
Process	Maps input to a predefined category.	Groups objects based on similarity.
Algorithm Examples	Decision Tree, Naive Bayes, SVM.	K-Means, Hierarchical Clustering.
Real-life Example	Sorting emails into "Inbox" or "Spam".	Grouping customers by similar buying habits.

2. Text Mining in Detail

Text Mining, also known as Text Data Mining or Knowledge Discovery from Text (KDT), is the process of extracting high-quality information from **unstructured text**.

While traditional data mining deals with structured data (databases/tables), text mining deals with documents, emails, social media posts, and articles.

Key Steps in Text Mining:

- ✓ Text Collection: Gathering documents from various sources like the web, file systems, or databases.
- ✓ Text Preprocessing: This is the most critical step. Since machines don't understand language like humans, the text must be cleaned:
- ✓ Tokenization: Breaking sentences into individual words (tokens).
- ✓ Stop-word Removal: Removing common words like "the," "is," "and" that don't carry much meaning.
- ✓ Stemming/Lemmatization: Reducing words to their root form (e.g., "running" becomes "run").
- ✓ Text Transformation (Feature Extraction): Converting text into a numerical format that a computer can process (e.g., TF-IDF or Word Embeddings).
- ✓ Pattern Discovery: Applying data mining algorithms (like classification or clustering) to find trends or associations within the text.
- ✓ Visualization: Representing the findings through word clouds, charts, or sentiment scores.

Applications of Text Mining:

- ✓ **Sentiment Analysis:** Determining if a movie review or tweet is positive, negative, or neutral.
- ✓ **Spam Filtering:** Identifying unwanted emails based on their content.
- ✓ **Fraud Detection:** Analyzing claims or documents to find suspicious patterns.
- ✓ **Information Retrieval:** Helping search engines like Google find the most relevant documents for your query.

© Ramit Pariyar

Why data mining is important? What are different aspects of security and privacy in data mining?

1. Why is Data Mining Important?

Data mining turns a "data swamp" into a "gold mine" of knowledge. Its importance can be summarized in three main points:

- ✓ Informed Decision Making: Instead of guessing, businesses use data mining to see what customers actually want. For example, Netflix uses it to decide which original shows to produce based on viewing habits.
- ✓ Predictive Analysis: It helps in forecasting future trends. Banks use it to predict which customers might default on loans, and retailers use it to predict which products will be in high demand during festivals.
- ✓ Efficiency & Optimization: It identifies bottlenecks. In manufacturing, data mining can reveal which machine is likely to fail soon, allowing for "preventative maintenance" before a costly breakdown occurs.

2. Security and Privacy in Data Mining

Since data mining often involves personal information (credit card numbers, health records, locations), security and privacy are major concerns. They are divided into several technical aspects:

A. Privacy-Preserving Data Mining (PPDM)

This is the most important aspect. It ensures that the "mining" process can happen without revealing the sensitive identity of individuals.

- ✓ **Anonymization:** Removing names, ID numbers, or specific addresses.
- ✓ **Data Perturbation:** Adding "noise" or slight random changes to the data so the individual values are wrong, but the overall patterns (the average or trend) remain accurate.

B. Data Integrity & Accuracy

Security isn't just about hiding data; it's about making sure the data isn't **tampered with**. If a hacker changes the data before it's mined, the resulting "knowledge" will be false and could lead to dangerous business or medical decisions.

© Ramit Pariyar

C. Access Control & Encryption

- ✓ **Access Control:** Only authorized users (like data scientists) should be able to run mining algorithms.
- ✓ **Encryption:** Data should be encrypted both while it is sitting in the Data Warehouse (at rest) and while it is being moved to the mining engine (in transit).

D. Social & Ethical Aspects

Even if the data is "legal" to mine, is it ethical?

- ✓ **Profiling:** Using data mining to discriminate against people based on race, religion, or gender (even if those fields are hidden, other data points might "reveal" them).
- ✓ **Consent:** Ensuring users know their data is being used for mining purposes and giving them the right to "opt-out."

Write short notes:

A. OLAP servers

OLAP (Online Analytical Processing) Servers are high-performance systems designed to support multi-dimensional data analysis. Unlike standard database servers that handle thousands of small transactions (OLTP), OLAP servers are optimized for complex "read-only" queries that summarize large volumes of data.

Depending on how they store and manage data, OLAP servers are classified into three main types:

1. ROLAP (Relational OLAP)

ROLAP servers sit on top of a traditional **Relational Database (RDBMS)**. They do not store a separate copy of the data; instead, they translate complex OLAP queries into standard SQL queries.

- ✓ **How it works:** It uses a "Star" or "Snowflake" schema to represent multi-dimensional data in flat tables.
- ✓ **Pros:** Highly scalable; can handle extremely large volumes of data.

- ✓ **Cons:** Slower performance compared to other types because SQL queries on huge tables can be time-consuming.

2. MOLAP (Multidimensional OLAP)

MOLAP servers store data in specialized **Multidimensional Arrays** (often called Data Cubes). They do not use tables.

- ✓ **How it works:** All calculations (aggregations) are pre-computed and stored in the cube before the user asks for them.
- ✓ **Pros:** Extremely fast response time because the answer is already "calculated" and ready.
- ✓ **Cons:** Not very scalable for massive data; requires a lot of storage space for the pre-computed cubes.

3. HOLAP (Hybrid OLAP)

As the name suggests, HOLAP is a combination of both ROLAP and MOLAP. It tries to offer the best of both worlds.

- ✓ **How it works:** It stores the large volume of detailed data in a relational database (ROLAP style) but keeps the pre-computed summary data (aggregations) in a multidimensional cube (MOLAP style).
- ✓ **Pros:** Balanced performance and scalability. Fast for summary queries and capable of handling large datasets.

B. Association rules

Association Rule Mining is one of the most popular data mining techniques. Its goal is to find interesting relationships or "associations" between variables in large datasets. It is most commonly used in Market Basket Analysis to find products that are frequently bought together.

1. The Concept

An association rule is usually written in the form of an implication:

X→Y This means: "If a customer buys item X, they are also likely to buy item Y."

Antecedent (X): The item found in the data.

Consequent (Y): The item found in combination with X.

2. The Three Core Measures

To decide if a rule is "strong" or "interesting," we use three mathematical metrics:

I. Support

It measures how frequently the itemset appears in the dataset.

$$\text{Support}(X \rightarrow Y) = \frac{\{\text{Number of transactions containing both } X \text{ and } Y\}}{\text{Total number of transactions}}$$

II. Confidence

It measures how often the rule has been found to be true. It is the conditional probability that a transaction containing X also contains Y.

$$\$\$\text{Confidence}(X \rightarrow Y) = \frac{\text{Support}(X \cup Y)}{\text{Support}(X)}$$

III. Lift

It measures the strength of a rule over the random co-occurrence of X and Y.

If **Lift > 1**, X and Y are positively correlated (buying X increases the chance of buying Y).

If **Lift = 1**, X and Y are independent.

3. Example: Grocery Store Analysis

Imagine a database of 100 customers:

40 customers buy **Bread**.

30 customers buy **Milk**.

20 customers buy **both Bread and Milk**.

Rule: $\{\text{Bread}\} \longrightarrow \{\text{Milk}\}$

Support = 20/100 = 20%

Confidence = 20/40 = 50% (Half of the bread buyers also bought milk).

© Ramit Pariyar

4. Famous Algorithm: Apriori

The **Apriori Algorithm** is the most widely used method for finding association rules. It works on the principle that:

"If an itemset is frequent, then all of its subsets must also be frequent."

This helps the system ignore infrequent combinations early on, saving a lot of computing power.

K-means clustering

K-means clustering is one of the most popular unsupervised machine learning algorithms used to group data points into a specific number of clusters (k). Since it is "unsupervised," the algorithm looks for patterns in the data without using pre-defined labels.

How the Algorithm Works

The process is iterative and follows these primary steps:

Initialization: You choose a value for k (the number of clusters) and randomly place k points in the data space. These are called **centroids**.

Assignment: Each data point is assigned to the nearest centroid based on a distance metric, typically **Euclidean distance**:

$$d(p,q)=\sqrt{\sum_{i=1}^{n}(1i-pi)^2}$$

Update: For each of the k clusters, the algorithm calculates the mean of all points assigned to it. The centroid is then moved to this new mean position.

Repeat: Steps 2 and 3 are repeated until the centroids no longer move significantly or a maximum number of iterations is reached.

Key Considerations

- ✧ Choosing k (The Elbow Method): Selecting the right number of clusters can be tricky. A common technique is the Elbow Method, where you plot the sum of squared distances (inertia) against different values of k. The "elbow" point of the

curve indicates the optimal k where adding more clusters doesn't significantly improve the fit.

- ✧ Sensitivity to Initialization: Because centroids start randomly, different starting positions can sometimes lead to different results. To fix this, a variation called K-means++ is often used to spread out the initial centroids more effectively.
- ✧ Feature Scaling: Since K-means relies on distance, it is vital to scale your data (e.g., using Min-Max scaling or Standardization). If one feature has a much larger range than others, it will dominate the distance calculation.

C. Support and confidence

In the context of Association Rule Mining (like the Apriori algorithm), Support and Confidence are the two primary metrics used to measure the strength and relevance of a relationship between items in a dataset.

Imagine we are looking at a set of transactions in a grocery store to see if buying Bread (A) leads to buying Butter (B).

1. Support

Support measures how **frequently** a specific item or a combination of items appears in the total dataset. It helps identify "frequent itemsets" so the algorithm doesn't waste time on rare combinations.

Formula:

$$\text{Support}(A \rightarrow B) = \frac{\text{Number of transactions containing both A and B}}{\text{Total number of transactions}}$$

High Support: The itemset appears often.

Low Support: The itemset is rare. If the support is below a certain "minimum support" threshold, the rule is usually discarded.

2. Confidence

Confidence measures how **reliable** the rule is. It tells us the probability that item B is purchased, given that item A has already been purchased.

© Ramit Pariyar

Formula:

$$\text{Confidence}(A \rightarrow B) = \frac{\text{Number of transactions containing both A and B}}{\text{Number of transactions containing A}}$$

Example: If Confidence is 0.80 (80%), it means that in 80 % of the cases where a customer bought Bread, they also bought Butter.

Limitation: High confidence can sometimes be misleading if item B is extremely popular overall (e.g., everyone buys Milk regardless of what else they buy).

End

© Ramit Pariyar